

Computational Economics Workshop

The University of Melbourne

John Stachurski

August 2024

Topics

Part 1: Workshop

- Background on scientific computing
- Current and future trends: AI-driven scientific computing
- Applications (hands on, using Python)

Part 2: Computational economics in action

- Joint with James Hansen and Yong Song

Key questions:

- What computational skills should economists learn in 2024?
- What are some interesting applications?

Please feel free to question / debate / share your experiences

Flow

- 1:00 – 2:30 Lecture 1
- 2:30 – 3:00 afternoon tea (staff lounge)
- 3:00 – 4:00 Lecture 2
- 4:00 – 4:15 break
- 4:15 – 5:00 Computational Economics in Action

Slides, code:

https://github.com/QuantEcon/melbourne_2024

Quick poll:

- Python programmers?
 - NumPy? Numba? PyTorch? JAX?
- Julia?
- MATLAB?
- C?
- Fortran?

Regular GPU users?

Prelude: AI-driven scientific computing

AI is changing the world

- image processing / computer vision
- speech recognition, translation
- scientific knowledge discovery
- forecasting and prediction
- generative AI

Plus killer drones, skynet, etc....

Prelude: AI-driven scientific computing

AI is changing the world

- image processing / computer vision
- speech recognition, translation
- scientific knowledge discovery
- forecasting and prediction
- generative AI

Plus killer drones, skynet, etc....

Projected spending on AI in 2024:

- Google: \$48 billion
- Microsoft: \$60 billion
- Meta: \$40 billion
- etc.

What kinds of problems are they solving?

Deep learning in two slides

Supervised deep learning: find a good approximation to an unknown functional relationship

$$y = f(x) \quad (x \in \mathbb{R}^d, y \in \mathbb{R})$$

Examples.

- x = sequence of words, y = next word
- x = weather sensor data, y = max temp tomorrow

Problem:

- observe $(x_i, y_i)_{i=1}^n$ and seek f such that $y_{n+1} \approx f(x_{n+1})$

Training: minimize the empirical loss

$$\ell(\theta) := \sum_{i=1}^n (y_i - f_{\theta}(x_i))^2 \quad \text{s.t.} \quad \theta \in \Theta$$

But what is $\{f_{\theta}\}_{\theta \in \Theta}$?

In the case of ANNs, we consider all f_{θ} having the form

$$f_{\theta} = \sigma \circ A_1 \circ \cdots \circ \sigma \circ A_{k-1} \circ \sigma \circ A_k$$

where

- $A_i x = W_i x + b_i$ is an affine map
- σ is a nonlinear “activation” function

Training: minimize the empirical loss

$$\ell(\theta) := \sum_{i=1}^n (y_i - f_{\theta}(x_i))^2 \quad \text{s.t.} \quad \theta \in \Theta$$

But what is $\{f_{\theta}\}_{\theta \in \Theta}$?

In the case of ANNs, we consider all f_{θ} having the form

$$f_{\theta} = \sigma \circ A_1 \circ \cdots \circ \sigma \circ A_{k-1} \circ \sigma \circ A_k$$

where

- $A_i x = W_i x + b_i$ is an affine map
- σ is a nonlinear “activation” function

Training: minimize the empirical loss

$$\ell(\theta) := \sum_{i=1}^n (y_i - f_{\theta}(x_i))^2 \quad \text{s.t.} \quad \theta \in \Theta$$

But what is $\{f_{\theta}\}_{\theta \in \Theta}$?

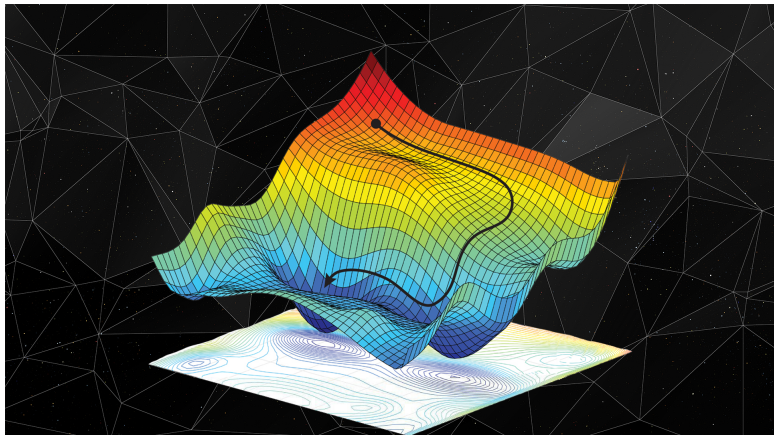
In the case of ANNs, we consider all f_{θ} having the form

$$f_{\theta} = \sigma \circ A_1 \circ \cdots \circ \sigma \circ A_{k-1} \circ \sigma \circ A_k$$

where

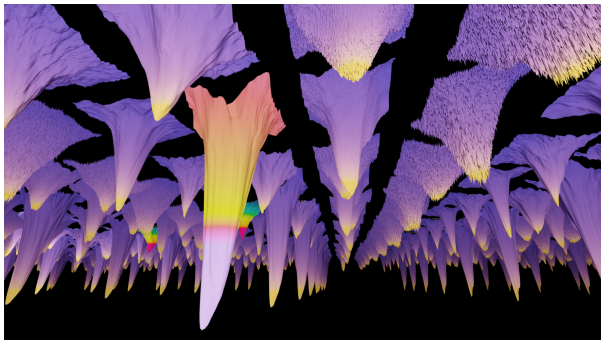
- $A_i x = W_i x + b_i$ is an affine map
- σ is a nonlinear “activation” function

Minimizing a smooth loss functions – what algorithm?



Source: <https://danielkhv.com/>

Deep learning: $\theta \in \mathbb{R}^d$ where $d = ?$



Source: <https://losslandscape.com/gallery/>

Hardware



“NVIDIA supercomputers are the factories of the AI industrial revolution.” – Jensen Huang

Software

Core elements

- automatic differentiation (for gradient descent)
- parallelization (GPUs! — how many?)
- Compilers / JIT-compilers

Crucially, these components must be well integrated

Popular platforms with these features

- Pytorch
- Google JAX (with Python)

```
import jax.numpy as jnp
from jax import grad, jit
```

```
def f( $\theta$ , x):
    for W, b in  $\theta$ :
        w = W @ x + b
        x = jnp.tanh(w)
    return x
```

```
def loss(theta, x, y):  
    return jnp.sum((y - f(theta, x))**2)
```

```
grad_loss = jit(grad(loss)) # Now use gradient descent
```

Source: JAX readthedocs

Summary

We can't afford the same hardware as OpenAI but we do have access to cheaper versions

And most of the relevant code is open source

Thus, AI-driven scientific computing is giving us many exciting and powerful new tools

- ▷ for deep learning and other kinds of mathematical modeling

Let's use them!